

Francielle Vargas
PhD in Computer Science - Artificial Intelligence
francielle.vargas@uchile.cl
<https://franciellevargas.github.io/>

Education	PhD in Computer Science and Computational Mathematics	2024
	University of São Paulo (USP)	
	MSc in Computer Science and Computational Mathematics	2017
	University of São Paulo (USP)	
	Software Engineering	2016
	Pontifical Catholic University of Minas Gerais (PUC)	
	BA in Linguistics	2014
	Federal University of Minas Gerais (UFMG)	

Research Projects	• 2026–Present. Principal Investigator. <i>Towards Trustworthy LLMs: Rationale Supervision and Evidence-Based Attention Alignment in RAG for Hallucination Mitigation.</i> Data and Artificial Intelligence Initiative, University of Chile, Chile. • 2026–Present. Research Collaborator. <i>Digital Trust Council Standards & Trustworthy AI Research.</i> Digital Trust Council & Harvard University, USA. • 2026–Present. Research Collaborator. <i>AI and Global Justice: Interdisciplinary Perspectives on Global Inequalities, and Inclusive AI Governance.</i> International Center for Global South Studies, University of Tübingen, Germany • 2024–Present. Principal Investigator. <i>Building Benchmarks for Hate Speech Detection with Moral Rationales.</i> University of Southern California, USA. • 2025. Postdoctoral Researcher. <i>Robust Augmented Retrieval for Natural Language Inference over Transformer-based Models.</i> São Paulo State University, Brazil & Idiap Research Institute, Switzerland. • 2024. Principal Investigator. <i>Responsible and Explainable Fact-Checking through Fine-Grained Factual Reasoning.</i> Google Latin America Research Award (LARA). • 2023. Principal Investigator. <i>Benchmarking Hate Speech Detection in Hausa Indigenous African Language.</i> University of São Paulo (USP), Brazil. • 2022. Research Collaborator. <i>Expanding Evaluation Data for Multilingual Protest News Detection.</i> Koç University, Turkey. • 2020. Doctoral Researcher. <i>Socially Responsible and Explainable Hate Speech Detection in Brazilian Portuguese.</i> University of São Paulo, Brazil.	
--------------------------	---	--

Awards & Honors	International	
	• 2026: Best PhD Dissertation Award in NLP, International Conference on Computational Processing of Portuguese and Galician (PROPOR). • 2025: Trevisan Prize for Students “AI for Good”, Bocconi University, Italy. • 2024: Google Latin America Research Award (LARA), Google Research. • 2024: ACL Diversity & Inclusion Award, Association for Computational Linguistics (EMNLP 2024 and NAACL 2024).	

National

- **2025:** Maria Carolina Monard Award for Best PhD Thesis in Artificial Intelligence, University of São Paulo, Brazil.
 - **2025:** Brazilian Computer Society (SBC) Thesis Award Finalist, Top 11 PhD Theses in Computer Science Nationwide.
 - **2025:** Brazilian Computer Society (SBC) Thesis Award Finalist – Multimedia, Hypermedia and Web, Top 6 PhD Theses Nationwide.
 - **2012, 2013:** Outstanding Academic Achievement Award with Honorable Mention, Federal University of Minas Gerais, Brazil.
-

Academic Experience	05/2026–Present Assistant Professor University of Chile, Chile Conducting research and teaching in Explainable Artificial Intelligence and Natural Language Processing as a member of the Data and Artificial Intelligence Initiative.
	01/2026–06/2026 Research Collaborator Digital Trust Council, Harvard University, USA Contributing to research on Trustworthy AI, and standards for digital trust.
	04/2025–04/2026 Postdoctoral Researcher São Paulo State University (Brazil) & Idiap Research Institute (Switzerland) Research on evidence-based retrieval, NLI, and interpretable Transformers models.
	01/2024–12/2024 Google LARA PhD Fellow Federal University of Minas Gerais, Brazil Recipient of the Google PhD Fellowship for research on Trustworthy AI, automated fact-checking, and language model interpretability.
	04/2024 Visiting PhD Student University of Southern California, USA Collaborated with the Morality and Language Lab on language model interpretability, fairness, and moral reasoning evaluation.
	11/2023 Invited Researcher Leibniz Institute for the Social Sciences, Germany Invited speaker at the Conference on Harmful Online Communication.
	08/2021–12/2021 Teaching Assistant – Neural Networks and Deep Learning University of São Paulo, Brazil Supported graduate-level teaching activities, student mentoring, and assessment.
	02/2020–06/2020 Teaching Assistant – Computing Theory and Compilers University of São Paulo, Brazil Supported undergraduate teaching, grading, and student supervision.
	07/2019–12/2023 PhD Researcher University of São Paulo, Brazil Conducted research on NLP, explainability, fairness, hate speech, and misinformation. Developed benchmark datasets and interpretable AI methods, resulting in publications in top-tier AI and NLP conferences and journals and award-winning research.

Publications

(*) Equal Contribution

Journal Papers

1. Francielle Vargas, Wolfgang Schmeisser-Nieto, Zohar Rabinovich, Thiago A.S. Pardo, Fabrício Benevenuto. Discourse Annotation Guideline for Low-Resource Languages. In *Natural Language Processing Journal*. Cambridge University Press, pp. 700-743. [pdf](#).
2. Francielle Vargas, Isabelle Carvalho, Thiago Pardo, Fabrício Benevenuto. Context-Aware and Expert Data Resources for Brazilian Portuguese Hate Speech Detection. In *Natural Language Processing Journal*. Cambridge University Press, pp. 435-456. [pdf](#)

Conference and Workshop Papers

1. Francielle Vargas, João Robiatti, Diego Alves, Lucas Pascotti Valem, Maximilian Seeth, Sebastián Ferrada, Ameeta Agrawal, Daniel Pedronette, André Freitas. Beyond Topical Similarity: Contrastive Evidence Retrieval with Interpretable Attention Alignment in RAG. *arXiv:2606.01482*. pp. 1-17. 2026. [pdf](#)
2. Francielle Vargas, Jackson Trager, Diego Alves, Surendrabikram Thapa, Matteo Guida, Berk Atil, Daryna Dementieva, Andrew Smart, Ameeta Agrawal. Self-Explaining Hate Speech Detection with Moral Rationales. *Findings of the Association for Computational Linguistics: ACL 2026*. pp. 1-18. San Diego, USA. [pdf](#)
3. Brage Eilertsen, Røskva Björgfinsdóttir, Francielle Vargas, Ali Ramezani-Kebrya. Aligning Attention with Human Rationales for Self-Explaining Hate Speech Detection. *Proceedings of the 40th Annual AAAI Conference on Artificial Intelligence*. pp. 1-15. Singapore, Singapore. [pdf](#)
4. Jackson Trager*, Francielle Vargas*, Diego Alves, Matteo Guida, Mikel Nguenajio, Ameeta Agrawal, Yalda Daryanai, Farzan Karimi-Malekabadi, Flor Miriam Plaza-del-Arco. MFTCXplain: A Multilingual Benchmark Dataset for Evaluating the Moral Reasoning of LLMs through Multi-hop Hate Speech Explanation. *Findings of the Association for Computational Linguistics: EMNLP 2025*. pp. 15709-15740, Suzhou, China. [pdf](#)
5. Isadora Salles, Francielle Vargas, Fabrício Benevenuto. HateBRXplain: A Benchmark Dataset with Human-Annotated Rationales for Explainable Hate Speech Detection in Brazilian Portuguese. *Proceedings of the 31st International Conference on Computational Linguistics (COLING 2025)*. pp. 6659-6669, Abu Dhabi, UAE. [pdf](#)
6. Agostina Calabrese, Christine de Kock, Debora Nozza, Flor Miriam Plaza-del-Arco, Zeerak Talat, Francielle Vargas. Proceedings of the 9th Workshop on Online Abuse and Harms (WOAH). *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (ACL 2025)*. pp. 1-549, Vienna, Austria. [pdf](#)
7. Francielle Vargas, Thiago Pardo, Fabrício Benevenuto. Socially Responsible and Explainable Automated Fact-Checking and Hate Speech Detection. *Anais Estendidos do XXXI Simpósio Brasileiro de Sistemas Multimídia e Web / Concurso de Teses e Dissertações*, Sociedade Brasileira de Computação. pp. 25-26, Porto Alegre, Brasil. [pdf](#)
8. Francielle Vargas, Thiago Pardo, Fabrício Benevenuto. Socially Responsible and Explainable Automated Fact-Checking and Hate Speech Detection. *Anais*

- do XXXVIII Concurso de Teses e Dissertações (CTD 2025), Sociedade Brasileira de Computação. pp. 75-84, Porto Alegre, Brasil. [pdf](#)
9. Francielle Vargas, Isadora Salles, Diego Alves, Ameeta Agrawal, Thiago Pardo, Fabrício Benevenuto. Improving Explainable Fact-Checking via Sentence-Level Factual Reasoning. *Proceedings of the 7th Fact Extraction and VERification Workshop (FEVER)*, co-located with EMNLP 2024. pp. 192–204, Miami, United States [pdf](#)
 10. Francielle Vargas, Samuel Guimarães, Shamsuddeen H. Muhammad, Diego Alves, Ibrahim Said Ahmad, Idris Abdulmumin, Diallo Mohamed, Thiago Pardo, Fabrício Benevenuto. HausaHate: An Expert Annotated Corpus for Hausa Hate Speech Detection. *Proceedings of the 8th Workshop on Online Abuse and Harms (WOAH)*, co-located with NAACL 2024. pp. 52–58. Mexico City, Mexico. [pdf](#)
 11. Surendrabikram Thapa, Kritesh Rauniyar, Farhan Jafri, Hariram Veeramani, Raghav Jain, Sandesh Jain, Francielle Vargas, Ali Hürriyetoglu, Usman Naseem. Extended Multimodal Hate Speech Event Detection During Russia-Ukraine Crisis - Shared Task at CASE 2024. *Proceedings of the 7th Workshop on Challenges and Applications of Automated Extraction of Socio-political Events from Text (CASE 2024)*, co-located with EACL 2024 pp. 221–228. St. Julians, Malta. [pdf](#)
 12. Francielle Vargas, Isabelle Carvalho, Ali Hürriyetoglu, Thiago A.S. Pardo, Fabrício Benevenuto (2023). Socially Responsible Hate Speech Detection: Can Classifiers Reflect Social Stereotypes?. *Proceedings of the 14th Recent Advances in Natural Language Processing (RANLP 2023)*. pp. 1187–1196. Varna, Bulgaria. [pdf](#)
 13. Francielle Vargas, Kokil Jaidka, Thiago A.S. Pardo, Fabrício Benevenuto (2023) Predicting Sentence-Level Factuality of News and Bias of Media Outlets. *Proceedings of the 14th Recent Advances in Natural Language Processing (RANLP 2023)*. pp. 1197–1206. Varna, Bulgaria. [pdf](#)
 14. Francielle Vargas, Isabelle Carvalho, Wolfgang Schmeisser-Nieto, Fabrício Benevenuto, Thiago A.S. Pardo. NoHateBrazil: A Brazilian Portuguese Text Offensiveness Analysis System. *Proceedings of the 14th Recent Advances in Natural Language Processing (RANLP 2023)*. pp.1180–1186. Varna, Bulgaria. [pdf](#)
 15. Surendrabikram Thapa, Farhan Jafri, Ali Hürriyetoglu, Francielle Vargas, Roy Ka-Wei Le, Usman Naseem. (2023) Multimodal Hate Speech Event Detection CASE Shared Task 4. *Proceedings of the 6th Workshop on Challenges and Applications of Automated Extraction of Socio-political Events from Text (CASE)*, co-located with RANLP 2023. pp.151-159. Varna, Bulgaria. [pdf](#)
 16. Francielle Vargas, Isabelle Carvalho, Fabiana R. Góes, Thiago A.S. Pardo, Fabrício Benevenuto. (2022) HateBR: A Large Expert Annotated Corpus of Brazilian Instagram Comments for Abusive Language Detection. *Proceedings of the 13th Conference on Language Resources and Evaluation (LREC 2022)*. pp. 7174–7183. Marseille, France. [pdf](#)
 17. Francielle Vargas, Jonas D’Alessandro, Zohar Rabinovich, Fabrício Benevenuto, Thiago A.S. Pardo. (2022) Rhetorical Structure Approach for Online Deception Detection: A Survey. *Proceedings of the 13th Conference on Language Resources and Evaluation (LREC 2022)*. pp. 5906-5915. Marseille, France. [pdf](#)
 18. Francielle Vargas, Thiago A. S. Pardo. Studying Dishonest Intentions in Brazilian Portuguese Texts. *Deceptive AI*. Springer: Communications in Computer and Information Science. vol 1296. pp. 166–178. [pdf](#)

19. Ali Hürriyetoğlu, Osman Mutlu, Fırat Duruşan, Onur Uca, Alaeddin Gürel, Benjamin J. Radford, Yaoyao Dai, Hansi Hettiarachchi, Niklas Stoehr, Tadashi Nomoto, Milena Slavcheva, Francielle Vargas, Aaqib Javid, Aaqib Javid, Erdem Yörük. Extended Multilingual Protest News Detection - Shared Task 1, CASE 2021 and 2022. *Proceedings of the 5th Workshop on Challenges and Applications of Automated Extraction of Socio-political Events from Text (CASE)*, co-located with EMNLP 2022. pp. 223–228. Abu Dhabi, Arab Emirates. [pdf](#)
20. Francielle Vargas, Fabiana R. Góes, Isabelle Carvalho, Fabrício Benevenuto, Thiago A.S. Pardo. Contextual-Lexicon Approach for Abusive Language Detection. *Proceedings of the 13th Recent Advances in Natural Language Processing (RANLP 2021)*. pp. 1442-1451. Held Online. [pdf](#)
21. Francielle Vargas, Fabrício Benevenuto, Thiago A.S. Pardo (2021). Towards Discourse-Aware Models for Multilingual Fake News Detection. *Proceedings of the Student Research Workshop Associated with RANLP 2021*. pp. 210-218. Held Online. [pdf](#)
22. Mateus T. Machado, Thiago A. S. Pardo, Evandro E. S. Ruiz, Ariani Di Felippo, Francielle Vargas. Implicit Opinion Aspect Clues in Portuguese Texts: Analysis and Categorization. *Proceedings of the 15th International Conference on the Computational Processing of Portuguese (PROPOR 2021)*. pp. 68-78. Fortaleza, Brazil. [pdf](#)
23. Jason R.C. Nurse, Francielle Vargas, Naeemul Hassan, Dakuo Wang, Panagiotis Andriotis, Amira Ghenai, Kokil Jaidka, Eni Mustafaraj, Kenneth Joseph and Brooke Foucault Welles. Towards A Diverse, Inclusive, Accessible and Equitable AAAI International Conference on Web and Social Media (ICWSM). *Workshop Proceedings of the 15th International AAAI Conference on Web and Social Media (ICWSM 2021)*. pp. 01-12. Held Online. [pdf](#)
24. Francielle Vargas, Thiago A. S. Pardo. Linguistic Rules for Fine-Grained Opinion Extraction. *Workshop Proceedings of the 14th International AAAI Conference on Web and Social Media (SocialSens)*, co-located with ICWSM 2020. pp. 1-6. Held Online. [pdf](#)
25. Francielle Vargas, Rodolfo Sanches Saraiva Dos Santos, Pedro Regattieri Rocha. Identifying Fine-Grained Opinion and Classifying Polarity on Coronavirus Pandemic. *Proceedings of the 9th Brazilian Conference on Intelligent Systems (BRACIS 2020)*. pp. 511-520. Rio Grande, Brazil. [pdf](#)
26. Francielle Vargas, Thiago A. S. Pardo. Aspect Clustering Methods for Sentiment Analysis. *Proceedings of the 13th International Conference on the Computational Processing of Portuguese (PROPOR 2018)*. pp.365-374. Canela, Brazil. [pdf](#)
27. Thiago A. S. Pardo, Jorge Baptista, Magali S. Duran, Maria das Graças Nunes, Fernando Nóbrega, Sandra Aluísio, Ariani Di Felippo, Eloize Seno, Raphael R. Silva, Rafael Anchieta, Henrico Brum, Márcio Dias, Rafael Martins, Erick Maziero, Jackson Souza, Francielle Vargas. The Coreference Annotation of the CSTNews Corpus. *Proceedings of the 2nd Workshop on Evaluation of Human Language Technologies for Iberian Languages (IberEval)*, co-located with SE-PLN 2017. pp. 102-112. Murcia, Spain. [pdf](#)

Non-peer reviewed articles

1. Francielle Vargas, and Thiago A. S. Pardo. Empirical Study on Clustering and Hierarchical Organization of Aspects for Opinion Mining. *Technical Institute of Mathematics and Computer Science's Reports, University of São Paulo*. São Carlos, Brazil. no.418. 48 pp. 2017.[pdf](#)

Invited Talks & Keynotes

- **2026** – Roundtable Lead, Algorithmic Fairness Across Alignment Procedures and Agentic Systems (AFAA) Workshop, ICLR 2026, Rio de Janeiro, Brazil.
Topic: *Fairness via Interpretable Alignment in Language Models*.
- **2026** – Keynote Speaker, 16th Brazilian-German Frontiers of Science and Technology Symposium (BRAGFOST), Alexander von Humboldt Foundation Frontiers of Research Symposia, Brazil-Germany.
Talk: *Bridging Interpretability and Alignment for Trustworthy Large Language Models*.
- **2026** – Keynote Speaker, 2nd RAIA Conference (Network for the Advancement of Artificial Intelligence), Institute of Mathematical and Computer Sciences (ICMC), University of São Paulo, Brazil.
Talk: *Interpretable Attention Alignment with Moral Rationales for Hate Speech Detection*.
- **2025** – Invited Talk, Applied Social Media Lab Synthesizer & Open Showcase, Berkman Klein Center for Internet & Society, Harvard University.
Posters: *Brazil#WithoutHate: Self-Explaining and Moral-Aware AI for Hate Speech Detection*; *Factuality and Transparency Are All RAG Needs! Self-Explaining Contrastive Evidence Re-ranking for Retrieval-Augmented Generation*.
- **2025** – Invited Research Supervisor, Summer School on Explainable and Trustworthy AI, University of Oslo (UiO), Norway.
Mentored two international MSc students in Explainable and Trustworthy AI.
- **2024** – Invited Talk, Language and Morality Lab, University of Southern California (USC), USA.
Talk: *Fighting Misinformation and Radicalism: Socially Responsible and Explainable Fact-Checking and Hate Speech Detection*.
- **2023** – Keynote Speaker, Conference on Harmful Online Communication, GESIS Leibniz Institute for the Social Sciences, Germany.
Talk: *Countering Harmful Online Communication in Brazil: Predicting Fine-Grained Factuality of News and Offensive Context of Social Media Comments*.

Organizing Committee

- **Lead Co-Organizer**, DeepXplain: Special Session on Explainable Deep Neural Networks for Responsible AI, *IJCNN 2025*, Rome, Italy.
- **Co-Organizer**, 9th Workshop on Online Abuse and Harms (WOAH), *ACL 2025*, Vienna, Austria.
- **Datasets Track Co-Chair**, *ICWSM 2023*.
- **Diversity, Inclusion, & Accessibility Chair**, *ICWSM 2021, 2022*.

Program Committee

1. **Area Chairing**: I have served as an Area Chair for leading c conferences, including:
 - International Joint Conference on Neural Networks (IJCNN).
Interpretable and Explainable AI Track
2. **Journal Reviewer**: I am a reviewer for some of the most prestigious international journals in Natural Language Processing (NLP), including:

- Natural Language Processing.
- Language Resources and Evaluation.
- Online Social Networks and Media.
- Expert Systems with Applications

3. **Conference Reviewer:** I am a reviewer for the main international conferences in Artificial Intelligence (AI) and Natural Language Processing (NLP), including:

- Annual Conference on Neural Information Processing Systems (NeurIPS)
- Annual Meeting of the Association for Computational Linguistics (ACL)
- Empirical Methods in Natural Language Processing (EMNLP)
- Annual Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics (NAACL)
- International Conference on Language Resources and Evaluation (LREC)
- International Conference on Computational Linguistics (COLING)
- International AAAI Conference on Web and Social Media (ICWSM)
- Conference on Information and Knowledge Management (CIKM)

Conference Participation / Presentation

- 2026: ACL, ICLR, PROPOR
 - 2025: EMNLP, CSBC
 - 2024: EMNLP, NAACL
 - 2023: ACL, RANLP
 - 2022: LREC
 - 2021: RANLP, ICWSM
 - 2020: ACL, EMNLP, LREC
-

Mentoring

- **PhD Student in Computer Science, Federal Fluminense University (UFF), Brazil (2026–present):** Co-supervisor of Juan David Nieto García, working on mechanistic interpretability and hate speech detection.
 - **MSc in Computer Science, Federal University of Minas Gerais (UFMG), Brazil (2024):** Co-supervised Isadora Salles in the development of *HateBRX-plain*, the first benchmark dataset for explainable hate speech detection in Brazilian Portuguese, published at COLING 2025.
 - **BSc in Computer Science, University of São Paulo (USP), Brazil (2020):** Co-supervised Lucas Sobral Fontes Cardado on a final project on opinion extraction and clustering from web consumer reviews, resulting in a Portuguese opinion mining framework.
-

Research Artifacts

Benchmark Datasets

- **HateBRXplain and HateBRMoralXplain:** Benchmark datasets with hate and moral human-annotated rationales for explainable hate speech detection.
- **MFTCXplain:** A multilingual benchmark dataset for evaluating the moral reasoning of large language models.
- **HateBR:** A benchmark dataset for hate speech detection in Brazilian Portuguese.
- **FactNews:** A benchmark dataset for sentence-level factuality prediction.
- **HausaHate:** A benchmark expert dataset for Hausa hate speech detection.
- **MOL:** A context-aware multilingual offensive lexicon.

Computational Methods

- **SMRA:** Supervised moral rational attention: Attention regularization with moral rationales for self-explaining hate speech detection.
- **CERA:** Contrastive evidence rationale attention via evidential inductive bias and interpretable attention alignment in RAG.
- **SRA:** Supervised rational attention for self-explaining hate speech detection.
- **SSA:** Post-hoc counter-stereotype explanations for bias assessment in hate speech detection models.
- **B+M:** Contextual BoW with human rationales and feature optimization for explainable hate speech detection.
- **SELFAR:** Sentence-level factual reasoning for explainable fact-checking.

Software Systems

- **FACTual:** A fact-checking and news source reliability estimation system.
- **NoHateBrazil:** A Brazilian Portuguese text offensiveness analysis system.
- **OPCluster:** A system for the extraction and clustering of opinions.

Advanced Training

- **2024:** 2nd Mexican NLP Summer School, NAACL 2024, Mexico.
- **2023:** 2nd Summer School on Deep Learning in NLP, RANLP 2023, Bulgaria.
- **2023:** IEEE Spoken Language Technology Workshop Hackathon, Qatar.
- **2022:** 2nd Advanced NLP School (ALPS), Université Grenoble Alpes, France.
- **2021:** 4th Advanced School in Big Data Analysis, ICMC-USP, Brazil.
- **2020:** 10th Lisbon Machine Learning School (LxMLS), INESC, Portugal.
- **2020:** Hackathon on Antisemitism in Social Media, Indiana University, USA.

Technical Skills

Python, NLTK, spaCy, Keras, Tensorflow, PyTorch, SQL, RDF(s), OWL, Unix.

Languages

English: Professional. **Spanish:** Basic. **Portuguese:** Native.

References

Dr. Ameeta Agrawal. Assistant Professor of the Department of Computer Science, Portland State University, USA.

Email: ameeta@pdx.edu.

I have established a successful and productive international collaboration with Dr. Agrawal, which has resulted in papers published to EMNLP 2024 and 2025. Our research addresses key topics such as multilingual natural language processing, explainability and interpretability, fairness, and moral reasoning evaluation of LLMs making contributions to the advancement of the field.

Dr. Roseli Romero. Associate Professor of Computer Science at the University of São Paulo, Brazil.

Email: rafrance@icmc.usp.br.

I had the privilege of working as a teaching assistant under the supervision of Dr. Roseli Romero for the Neural Networks and Deep Learning course in 2021. Additionally, Dr. Romero and I, in collaboration with Dr. Jackson Trager from the University of Southern California, organized the special session Explainable Deep Neural Networks for Responsible AI (DeepXplain), co-located with International Joint Conference on Neural Networks (IJCNN 2025) in Italy. The special session aims to foster advancements in interpretability and the trustworthy use of deep neural networks.

Dr. Debora Nozza. Assistant Professor in the Computer Science Department at Bocconi University, Italy.

Email: debora.nozza@unibocconi.it

I first met Dr. Debora Nozza in person at EMNLP 2025 in Miami. We later worked together in the ACL community as members of the organizing committee of the 9th Workshop on Online Abuse and Harms (WOAH), co-located with the ACL 2025, which she co-chaired. Our interaction was focused on workshop organization and program-related activities, reflecting a shared interest in hate speech research and socially-aware and responsible NLP. I am currently collaborating with her postdoctoral researcher, Dr. Arianna Muti, on a project I am leading focused on language model interpretability, hate speech, and multicultural contexts. Debora has also been exceptionally generous and supportive of my career development, including endorsing my application for the ACL Dissertation Award in 2026.

Dra. Flor Plaza-del-Arco. Assistant Professor Institute of Advanced Computer Science, Leiden University, Netherlands.

Email: f.m.plaza.del.arco@liacs.leidenuniv.nl

I have had the pleasure of working with Dra. Flor Plaza-del-Arco to organize the 9th Workshop on Online Abuse and Harms (WOAH), co-located with the 63rd Annual Meeting of the Association for Computational Linguistics (ACL 2025). Together with other outstanding researchers, we worked to bring attention to the critical issues of hate speech and misinformation, aiming to foster a space for sharing insights and developing solutions. Additionally, Dra. Flor Plaza-del-Arco is a co-author of our paper "MFTCXplain: A Multilingual Benchmark Dataset for Evaluating the Moral Reasoning of LLMs through Multi-hop Hate Speech Explanations, published in EMNLP 2025.

Dr. Fabricio Benevenuto. Associate Professor of Computer Science at Federal University of Minas Gerais, Brazil.

Email: fabricio@dcc.ufmg.br.

I have had the privilege of working closely with Dr. Fabrício Benevenuto, who was my co-advisor during my PhD. Dr. Benevenuto is an internationally recognized expert in the fields of disinformation and hate speech, consistently ranked among the most influential researchers in the world. As a result of our collaboration, we have published

12 papers in top-tier NLP conferences, and I also had the opportunity to co-advise his master’s student, which resulted in an international publication at COLING 2025.

Dr. Morteza Dehghani. Director of the Center for Computational Language Sciences and Professor of Psychology and Computer Science at University of Southern California, USA.

Email: mdehghan@usc.edu.

I had the incredible opportunity to visit the MOLA – Morality and Language Lab at the University of Southern California (USC), coordinated by Dr. Morteza Dehghani. My visit focused on interdisciplinary research at the intersection of computational linguistics, psychology, and explainable and responsible AI. Dr. Dehghani’s expertise in cognitive modeling, computational social science, and natural language processing greatly enriched my understanding of how social-psychological theories can inform computational models, particularly in the context of Large Language Models (LLMs). As a result of this visit, Dr. Dehghani’s students and I published a paper to EMNLP 2025, in which we introduce the first multilingual benchmark dataset for evaluating the moral reasoning capabilities of LLMs.

Dr. Eduard Hovy. Executive Director of Melbourne University, Melbourne, Australia, and Associate Professor in the Language Technologies Institute at Carnegie Mellon University, USA.

Emails: ehovy@andrew.cmu.edu and hovy@cmu.edu.

During my Ph.D., I had the privilege of receiving guidance from Dr. Eduardo Hovy, who generously shared valuable insights on topics related to my research. We had the opportunity to meet in person at RANLP 2023 in Bulgaria, where we discussed key aspects relevant to the advancement of my work. I am currently collaborating with his Ph.D. student, Matteo Guida from the University of Melbourne. Our collaboration has resulted in a paper published at EMNLP 2025, a second work currently under submission to ACL 2026, and ongoing joint work toward a journal submission to TACL, focused on monolingual and multilingual expert-annotated benchmarks and self-explaining methods for hate speech and moral reasoning in large language models.